


Proximity-based Outlier Detection

- Objects far away from the others are outliers
- The proximity of an outlier deviates significantly from that of most of the others in the data set
- Distance-based outlier detection: An object o is an outlier if its neighborhood does not have enough other points
- Density-based outlier detection: An object o is an outlier if its density is relatively much lower than that of its neighbors

Depth-based Methods

- Organize data objects in layers with various depths
 - The shallow layers are more likely to contain outliers
- Example: Peeling, Depth contours
- Complexity $O(N^{\lceil k/2 \rceil})$ for k-d datasets
 - Unacceptable for $k > 2$

Depth-based Outliers: Example



Distance-based Outliers

- A $DB(p, D)$ -outlier is an object O in a dataset T such that at least a fraction p of the objects in T lie at a distance greater than distance D from O
- The larger D , the more outlying
- The larger p , the more outlying

Index-based Algorithms

- Find DB(p , D) outliers in T with n objects
 - Find an objects having at most $\lfloor n(1-p) \rfloor$ neighbors with radius D
- Algorithm
 - Build a standard multidimensional index
 - Search every object O with radius D
 - If there are at least $\lfloor n(1-p) \rfloor$ neighbors, O is not an outlier
 - Else, output O

Index-based Algorithms: Pros & Cons

- Complexity of search $O(kN^2)$
 - More scalable with dimensionality than depth-based approaches
- Building a right index is very costly
 - Index building cost renders the index-based algorithms non-competitive

A Naïve Nested-loop Algorithm

- For $j=1$ to n do
 - Set $\text{count}_j=0$;
 - For $k=1$ to n do if $(\text{dist}(j,k)<D)$ then count_j++ ;
 - If $\text{count}_j \leq \lfloor n(1-p) \rfloor$ then output j as an outlier;
- No explicit index construction
 - $O(N^2)$
- Many database scans

Improving Nested-loop Algorithm


- Once an object has at least $\lfloor n(1-p) \rfloor$ neighbors with radius D , no need to count further
- Use the data in main memory as much as possible
 - Reduce the number of database scans

Block-based Nested-loop Algorithm


- Partition the available memory into two blocks with an equivalent size
- Fill the first block, compare objects in the block, mark non-outliers
- Read remaining objects into the second block, compare objects from the first and second block
 - Mark non-outliers, only compare potential outliers in the first block
 - Output unmarked objects in the first block as outliers
- Swap the names of the first and second blocks, until all objects have been processed

Example


Dataset has four blocks: A, B, C, and D



Compare objects in A (1 read)




Compare objects in A to those in B, C, and D (3 reads)





Compare objects in D (0 read)






Compare objects in D to those in A (0 read)





Compare objects in D to those in B and C (2 reads)



Compare objects in C to those in C, D, A, and B (2 reads)



Compare objects in B to those in B, C, A, and D (2 reads)



10 blocks are read in total
 $10/4=2.5$ passes over T


Nested-loop Algorithm: Analysis

- The data set is partition into n blocks
- Total number of block reads:
 - $n + (n-2)(n-1) = n^2 - 2n + 2$
- The number of passes over the dataset
 - $\geq (n-2)$
- Many passes for large datasets

A Cell-based Approach

$$L_1(C_{x,y}) = \{C_{u,v} \mid |u - x| \leq 1, |v - y| \leq 1, C_{u,v} \neq C_{x,y}\}$$

$$l = \frac{D}{2\sqrt{2}} \quad L_2(C_{x,y}) = \{C_{u,v} \mid |u - x| \leq 3, |v - y| \leq 3, C_{u,v} \notin L_1(C_{x,y}), C_{u,v} \neq C_{x,y}\}$$



M+ objects in $C_{x,y}$
 → no outlier in $C_{x,y}$

M+ objects in $C_{x,y} \cup L_1(C_{x,y})$
 → no outlier in $C_{x,y}$

M- objects in $C_{x,y} \cup L_1(C_{x,y}) \cup L_2(C_{x,y})$
 → all objects in $C_{x,y}$ are outliers

The Algorithm

- Quantize each object to its appropriate cell
- Label all cells having $m+$ objects red
 - No outlier in red cells
- Label L_1 neighbours of red cells, and cells having $m+$ objects in $C_{x,y} \cup L_1(C_{x,y})$ pink
 - No outlier in pink cells
- Output objects in cells having $m-$ objects in $C_{x,y} \cup L_1(C_{x,y}) \cup L_2(C_{x,y})$ as outliers
- For remaining cells, check them one by one

Cell-based Approach: Analysis

- A typical cell has 8 L_1 neighbours and 40 L_2 neighbours
- Complexity: $O(m+N)$ (m : # of cells)
 - The worst case: no red/pink cell at all
 - In practice, many red/pink cells
- The method can be easily generalized to k -d space and other distance functions

Handling Large Datasets


- Where do we need page reads?
 - Quantize objects to cells: 1 pass
 - Object-pairwise: many passes
- Idea: only keep white objects in main memory
 - White objects are in cells not red nor pink

Reducing Disk Reads

- Classify pages in datasets
 - A: contain some white objects
 - B: contain no white objects but L_2 neighbours of white objects
 - C: other pages
 - Object-pairwise don't need class C pages
- Scheduling pages A and B properly
- At most 3 passes

Density-based Local Outlier

Both o1 and o2 are outliers
Distance-based methods
can detect o1, but not o2



Intuition

- Outliers comparing to their local neighborhoods, instead of the global data distribution
- The density around an outlier object is significantly different from the density around its neighbors
- Use the relative density of an object against its neighbors as the indicator of the degree of the object being outliers

K-Distance

- The k-distance of p is the distance between p and its k-th nearest neighbor
- In a set D of points, for any positive integer k, the k-distance of object p, denoted as k-distance(p), is the distance $d(p, o)$ between p and an object o such that
 - For at least k objects $o' \in D \setminus \{p\}$, $d(p, o') \leq d(p, o)$
 - For at most (k-1) objects $o' \in D \setminus \{p\}$, $d(p, o') < d(p, o)$


K-distance Neighborhood

- Given the k-distance of p , the k-distance neighborhood of p contains every object whose distance from p is not greater than the k-distance
 - $N_{k\text{-distance}(p)}(p) = \{q \in D \setminus \{p\} \mid d(p, q) \leq k\text{-distance}(p)\}$
 - $N_{k\text{-distance}(p)}(p)$ can be written as $N_k(p)$

Reachability Distance

- The reachability distance of object p with respect to object o is $\text{reach-dist}_k(p, o) = \max\{k\text{-distance}(o), d(p, o)\}$

If p and o are close to each other, $\text{reach-dist}(p, o)$ is the k -distance, otherwise, it is the real distance




Local Reachability Density

$$lrd_k(o) = \frac{|N_k(o)|}{\sum_{o' \in N_k(o)} reachdist_k(o' \leftarrow o)}$$


Local outlier factor

$$LOF_{MinPts}(p) = \frac{\sum_{o \in N_{MinPts}(p)} \frac{lrd_{MinPts}(o)}{lrd_{MinPts}(p)}}{|N_{MinPts}(p)|}$$



Examples



Example dataset




Examples



Clustering-based Outlier Detection

- An object is an outlier if
 - It does not belong to any cluster;
 - There is a large distance between the object and its closest cluster ; or
 - It belongs to a small or sparse cluster




Classification-based Outlier Detection

- Train a classification model that can distinguish “normal” data from outliers
- A brute-force approach: Consider a training set that contains some samples labeled as “normal” and others labeled as “outlier”
 - A training set in practice is typically heavily biased: the number of “normal” samples likely far exceeds that of outlier samples
 - Cannot detect unseen anomaly

One-Class Model

- A classifier is built to describe only the normal class
- Learn the decision boundary of the normal class using classification methods such as SVM
- Any samples that do not belong to the normal class (not within the decision boundary) are declared as outliers
- Advantage: can detect new outliers that may not appear close to any outlier objects in the training set
- Extension: Normal objects may belong to multiple classes


One-Class Model



Semi-Supervised Learning Methods

- Combine classification-based and clustering-based methods
- Method
 - Use a clustering-based approach to find a large cluster, C , and a small cluster, C_1
 - Since some objects in C carry the label “normal”, treat all objects in C as normal
 - Use the one-class model of this cluster to identify normal objects in outlier detection
 - Since some objects in cluster C_1 carry the label “outlier”, declare all objects in C_1 as outliers
 - Any object that does not fall into the model for C (such as a) is considered an outlier as well

Example



- objects with label "normal"
- objects with label "outlier"
- objects without label

Pros and Cons

- Pros: Outlier detection is fast
- Cons: Quality heavily depends on the availability and quality of the training set,
 - It is often difficult to obtain representative and high-quality training data

Contextual Outliers

- An outlier object deviates significantly based on a selected context
 - Ex. Is 10C in Vancouver an outlier? (depending on summer or winter?)
- Attributes of data objects should be divided into two groups
 - Contextual attributes: defines the context, e.g., time & location
 - Behavioral attributes: characteristics of the object, used in outlier evaluation, e.g., temperature
- A generalization of local outliers—whose density significantly deviates from its local area
- Challenge: how to define or formulate meaningful context?

Detection of Contextual Outliers

- If the contexts can be clearly identified, transform it to conventional outlier detection
 - Identify the context of the object using the contextual attributes
 - Calculate the outlier score for the object in the context using a conventional outlier detection method


Example

- Detect outlier customers in the context of customer groups
 - Contextual attributes: age group, postal code
 - Behavioral attributes: the number of transactions per year, annual total transaction amount
- Method
 - Locate c 's context;
 - Compare c with the other customers in the same group; and
 - Use a conventional outlier detection method

Modeling Normal Behavior

- Model the “normal” behavior with respect to contexts
 - Use a training data set to train a model that predicts the expected behavior attribute values with respect to the contextual attribute values
 - An object is a contextual outlier if its behavior attribute values significantly deviate from the values predicted by the model
- Use a prediction model to link the contexts and behavior
 - Avoid explicit identification of specific contexts
 - Some possible methods: regression, Markov Models, and Finite State Automaton ...

Collective Outliers




- Objects as a group deviate significantly from the entire data
- Examine the structure of the data set, i.e, the relationships between multiple data objects
 - The structures are often not explicitly defined, and have to be discovered as part of the outlier detection process.

Detecting High Dimensional Outliers

- Interpretability of outliers
 - Which subspaces manifest the outliers or an assessment regarding the “outlying-ness” of the objects
- Data sparsity: data in high-D spaces are often sparse
 - The distance between objects becomes heavily dominated by noise as the dimensionality increases
- Data subspaces
 - Local behavior and patterns of data
- Scalability with respect to dimensionality
 - The number of subspaces increases exponentially

Angle-based Outliers



To-Do List

- Read the rest of Chapter 12.4